Towards Network-level Efficiency for Cloud Storage Services

Zhenhua Li
Tsinghua University
Peking University
lizhenhua1983@tsinghua.edu.cn

Cheng Jin
University of Minnesota
Twin Cities
cheng@cs.umn.edu

Tianyin Xu
University of California
San Diego
tixu@cs.ucsd.edu

Christo Wilson
Northeastern University
Boston, MA, US
cbw@ccs.neu.edu

Yao Liu
State University of New York
Binghamton University
yaoliu@cs.binghamton.edu

Linsong Cheng
Tsinghua University
Beijing, China
chengls10@mails.tsinghua.edu.cn

Yunhao Liu
Tsinghua University
Beijing, China
yunhao@tsinghua.edu.cn

Yafei Dai
Peking University
Beijing, China
dyf@pku.edu.cn

Zhi-Li Zhang
University of Minnesota
Twin Cities
zhzhang@cs.umn.edu

ABSTRACT

Cloud storage services such as Dropbox, Google Drive, and Microsoft OneDrive provide users with a convenient and reliable way to store and share data from anywhere, on any device, and at any time. The cornerstone of these services is the *data synchronization* (sync) operation which *automatically* maps the changes in users' local filesystems to the cloud via a series of network communications in a *timely* manner. If not designed properly, however, the tremendous amount of data sync traffic can potentially cause (financial) pains to both service providers and users.

This paper addresses a simple yet critical question: *Is the current data sync traffic of cloud storage services efficiently used?* We first define a novel metric named *TUE* to quantify the *Traffic Usage Efficiency* of data synchronization. Based on both real-world traces and comprehensive experiments, we study and characterize the *TUE* of six widely used cloud storage services. Our results demonstrate that a considerable portion of the data sync traffic is in a sense wasteful, and can be effectively avoided or significantly reduced via carefully designed data sync mechanisms. All in all, our study of *TUE* of cloud storage services not only provides guidance for service providers to develop more efficient, traffic-economic services, but also helps users pick appropriate services that best fit their needs and budgets.

Categories and Subject Descriptors

C.2.4 [**Computer-communication Networks**]: Distributed Systems—*Distributed applications*; D.4.3 [**Operating Systems**]: File Systems Management—*Distributed file systems*

General Terms

Design, Experimentation, Measurement, Performance

Keywords

Cloud storage service; network-level efficiency; data synchronization; traffic usage efficiency

1. INTRODUCTION

Cloud storage services such as Dropbox, Google Drive, and Microsoft OneDrive (renamed from SkyDrive since Feb. 2014) provide users with a convenient and reliable way to store and share data from anywhere, on any device, and at any time. The users' data (*e.g.,* documents, photos, and music) stored in cloud storage are *automatically* synchronized across all the designated devices (*e.g.,* PCs, tablets, and smartphones) connected to the cloud in a *timely* manner. With multiplicity of devices – especially mobile devices – that users possess today, such "anywhere, anytime" features significantly simplify data management and consistency maintenance, and thus provide an ideal tool for data sharing and collaboration.

In a few short years, cloud storage services have reached phenomenal levels of success, with the user base growing rapidly. For example, Microsoft OneDrive claims that over 200 million customers have stored more than 14 PB of data using their service [9], while Dropbox has claimed more than 100 million users who store or update *1 billion* files every day [6]. Despite the late entry into this market (in Apr. 2012), Google Drive obtained 10 million users just in its first two months [7].

The key operation of cloud storage services is *data synchronization* (sync) which automatically maps the changes in users' local filesystems to the cloud via a series of network communications. Figure 1 demonstrates the general data sync principle. In a cloud storage service, the user usually needs to assign a designated local folder (called a "sync folder") in which every file operation is noticed and synchronized to the cloud by the client software developed by the service provider. Synchronizing a file involves a sequence of data sync events, such as transferring the data index, data content, sync notification, sync status/statistics, and sync acknowledgement. Naturally, each data sync event incurs network traffic. In this paper, this traffic is referred to as *data sync traffic*.

Figure 1: Data synchronization principle.

If not designed properly, the amount of data sync traffic can potentially cause (financial) pains to both providers and users of cloud storage services. From the providers' perspective, the aggregate sync traffic from all users is enormous (given the huge number of files uploaded and modified each day!). This imposes a heavy burden in terms of infrastructure support and monetary costs (*e.g.,* as payments to ISPs or cloud infrastructure providers). To get a quantitative understanding, we analyze a recent large-scale Dropbox trace [12] collected at the ISP level [25]. The analysis reveals: (1) The sync traffic contributes to more than 90% of the total service traffic. Note that the total service traffic is equivalent to one third of the traffic consumed by YouTube [25]; (2) Data synchronization of a file (sometimes a batch of files) generates 2.8 MB of *inbound* (client to cloud) traffic and 5.18 MB of *outbound* (cloud to client) traffic on average. According to the Amazon S3 pricing policy [1] (Dropbox stores all the data content in S3 and S3 only charges for outbound traffic), the Dropbox traffic would consume nearly $0.05/GB × *5.18 MB × 1 billion* = $260,000 every day[1]. These costs grow even further when we consider that *all* cloud storage service providers must bear similar costs, not just Dropbox [4].

Data sync traffic can also bring considerable (and unexpected) financial costs to end users, despite that basic cloud storage services are generally free. News media has reported about user complaints of unexpected, additional charges from ISPs, typically from mobile users with limited data usage caps [8, 2]. As a consequence, some users have warned: *"Keep a close eye on your data usage if you have a mobile cloud storage app."* In addition, some cloud storage applications (*e.g.,* large data backup [3]) are also impaired by the bandwidth constraints between the user clients and the cloud. This limitation is regarded as the "dirty secrets" of cloud storage services [5]. Hence users likewise would also benefit from more efficient sync traffic usage.

This paper addresses a simple yet critical question: *Is the current data sync traffic of cloud storage services efficiently used?* Our goal is to quantify and optimize the *efficiency* of data sync traffic usage, *i.e.,* the pivotal network-level efficiency for cloud storage services. Without impairing user experience, providers would like to limit data sync traffic as much as possible to reduce operational costs. On the other side, users also desire more efficient traffic usage, which can save money and result in better quality of experience. Although several studies have measured cloud storage services [30, 33, 25, 20, 24, 36, 32, 48], none have addressed the issue of sync traffic efficiency using real-world, large-scale data from multiple cloud storage services.

To answer the question thoroughly, we first define a novel metric named **TUE** to quantify the *Traffic Usage Efficiency* of data synchronization. Borrowing a term similar to *PUE* (*i.e.,* the *Power Usage Effectiveness* = $\frac{\text{Total facility power}}{\text{IT equipment power}}$ [14], a widely adopted metric for evaluating the cloud computing energy efficiency), we define

$$TUE = \frac{\text{Total data sync traffic}}{\text{Data update size}}. \qquad (1)$$

When a file is updated (*e.g.,* created, modified, or deleted) at the user side, the *data update size* denotes the size of altered bits relative to the cloud-stored file [2]. From the users' point of view, the data update size is an intuitive and natural signifier about how much traffic *should* be consumed. Compared with the absolute value of sync traffic (used in previous studies), *TUE* better reveals the essential traffic harnessing capability of cloud storage services.

In order to gain a practical and in-depth understanding of *TUE*, we collect a real-world user trace and conduct comprehensive benchmark experiments of six widely used cloud storage services, including Google Drive, OneDrive, Dropbox, Box, Ubuntu One, and SugarSync. We examine key *impact factors* and *design choices* that are common across all of these services. *Impact factors* include file size, file operation, data update size, network environment, hardware configuration, access method, and so on. Here the "access method" refers to PC client software, web browsers, and mobile apps. *Design choices* (of data sync mechanisms) include data sync granularity, data compression level, data deduplication granularity, and sync deferment (for improved batching).

By analyzing these factors and choices, we are able to thoroughly unravel the *TUE* related characteristics, design tradeoffs, and optimization opportunities of these state-of-the-art cloud storage services. The major findings in this paper and their implications are summarized as follows:

- The majority (77%) of files in our collected trace are small in size (less than 100 KB). Nearly two-thirds (66%) of these small files can be logically combined into larger files for *batched data sync* (BDS) in order to reduce sync traffic. However, only Dropbox and Ubuntu One have partially implemented BDS so far.

- The majority (84%) of files are modified by users at least once. Unfortunately, most of today's cloud storage services are built on top of RESTful infrastructure (e.g., Amazon S3, Microsoft Azure, and OpenStack Swift) that typically only support data access operations at the *full-file* level [26, 17]. For these services, enabling the efficient *incremental data sync* (IDS) mechanism requires an extra mid-layer for transforming MODIFY into GET + PUT + DELETE file operations. Given that file modifications frequently happen, implementing IDS is worthwhile for improved network-level efficiency.

- 52% of files can be effectively compressed and 18% of files can be deduplicated. Nevertheless, Google Drive, OneDrive, Box, and SugarSync never compress or deduplicate data. Even for Dropbox and Ubuntu One, the effect of compression and deduplication is largely influenced by the access method.

- Implementing compression and block-level deduplication together is technically challenging. Based on our trace analysis, we suggest providers to implement compression and full-file deduplication because the combination of these two techniques is sufficient to provide efficient usage of sync traffic.

---

[1]We assume that there is no special pricing contract between Dropbox and Amazon S3, so our calculation of the traffic costs may involve potential overestimation.

[2]If data compression is utilized by the cloud storage service, the *data update size* denotes the compressed size of altered bits.

- Frequent modifications to a file often lead to large *TUE*. For instance, for 8.5% of Dropbox users, more than 10% of their sync traffic is caused by frequent modifications [36]. Some services deal with this issue by batching file updates using a fixed sync deferment. However, fixed sync deferments are inefficient in some scenarios. We propose an *adaptive sync defer* (ASD) mechanism to overcome this limitation.

- In the presence of frequent file modifications, surprisingly, users with relatively "poor" hardware or Internet access save on sync traffic, because their file updates are naturally batched.

In a nutshell, our research findings demonstrate that for today's cloud storage services, a considerable portion of the data sync traffic is in a sense wasteful, and can be effectively avoided or significantly reduced through carefully designed data sync mechanisms. In other words, there is plenty of space for optimizing the network-level efficiency of these services. Our study of *TUE* provides guidance in two folds: (1) help service providers develop more efficient, traffic-economic cloud storage services; and (2) help end users select appropriate services that best fit their needs and budgets.

Roadmap. In the remainder of the paper, we first describe the common design framework of cloud storage services in § 2, and then introduce our research methodology in § 3. Next, we present research results and findings, broken down logically into three areas: simple file operations (§ 4), compression and deduplication (§ 5), and frequent file modifications (§ 6). Finally, we discuss the tradeoffs for cloud storage system design in § 7, review related work in § 8, and conclude the paper with future work in § 9.

2. COMMON DESIGN FRAMEWORK OF CLOUD STORAGE SERVICES

From the perspective of sync traffic usage, the common design framework of cloud storage services involves a number of *impact factors* and *design choices*, which can be on the client side, server (cloud) side, or network side. The *impact factors* refer to those (*objective*) factors such as the client location, hardware, file size, data update size, network environment, and so on that must be accounted for in the design and usage of cloud storage services. The *design choices* (of data sync mechanisms) refer to those (*subjective*) design decisions which the system designers make, such as the data sync granularity, data compression level, data deduplication granularity, and so forth.

Both the impact factors and design choices may influence the data sync *TUE*. To avoid being trapped by trivial or elusive issues, we select key impact factors and design choices according to the following two rules:

- *Rule 1: The impact factors should be relatively constant or stable, so that our research results can be easily repeated.*

- *Rule 2: The design choices should be measurable and service/implementation independent, so as to make our research methodology widely applicable.*

Following *Rule 1*, we do not study impact factors such as sync delay[3], cloud server location, etc. For example, we observe that uploading a 1-MB JPEG photo to Google Drive may incur an elusive sync delay varying between several seconds and several minutes (under different network environments). Instead, we choose to study the *sync traffic*, which is almost invariable in all cases.

---

[3]*Sync delay* measures how long the user client synchronizes a file to the cloud.

Table 1: Key impact factors and *design choices*.

Client side	Client location Client hardware
	Access method File size File operation
	Data update size Data update rate
	Data compression level Sync deferment
Server side	*Data sync granularity*
	Data deduplication granularity
	(*Data compression level*) [*]
Network side	Sync traffic Bandwidth Latency

[*] Note: The server-side data compression level may be different from the client-side data compression level.

Besides, we observe that the cloud server location serving a given file is not constant. This is because a cloud storage service usually hosts a user's files across multiple geographically dispersed data centers, and it often migrates or copies a file from one cloud server to another. Instead, we record the *bandwidth* and *delay* between the client and the cloud, as they can be reproduced using client-side methods (introduced in § 3.2).

Following *Rule 2*, we do not consider design choices such as the metadata structures, file segmentation and replication on the cloud side, because they require specific knowledge of the back-end cloud implementation. For example, the metadata structure (including the list of the user's files, their attributes, and indices to where the files can be found inside the cloud) cannot be extracted from the network communication packets, because almost all the commercial cloud storage services have encrypted their application-layer data in certain (unknown) ways.

In the end, ten key impact factors and four design choices are selected, as listed in Table 1. Some of them are self-explanatory or have been explained before. Below we further explain a few:

- *File operation* includes file creation, file deletion, file modification, and frequent file modifications.

- *Data update rate* denotes how often a file operation happens.

- *Sync deferment*. When frequent file modifications happen, some cloud storage services *intentionally* defer the sync process for a certain period of time for batching file updates.

- *Data sync granularity*. A file operation is synchronized to the cloud either in a *full-file* granularity or in an incremental, *chunk-level* granularity. When the former is adopted, the whole updated file is delivered to the cloud; when the latter is adopted, only those file chunks that contain altered bits (relative to the file stored in the cloud) are delivered.

- *Data deduplication granularity* denotes the *unit* at which data fingerprints are computed and compared to avoid delivering duplicate data units to the cloud. The unit can be either a full file or a file block. Note that data deduplication can be performed across different files owned by different users.

- *Bandwidth* is defined as the peak upload rate between the client and the cloud server. We measure it by uploading a large file to the cloud and meanwhile recording the network traffic with the Wireshark network protocol analyzer [18].

- *Latency* is defined as the round trip time (RTT) between the client and the cloud. We measure it by using the standard Ping command.

Table 2: Number of users and files recorded in our collected cloud storage trace.

	Google Drive	OneDrive	Dropbox	Box	Ubuntu One	SugarSync
Number of users	33	24	55	13	13	15
Number of files	32677	17903	106493	19995	27281	18283

Table 3: File attributes recorded in our collected trace.

| User name	File name MD5	Original file size
Compressed file size	Creation time	Last modification time
Full-file MD5	128 KB/256 KB/512 KB/1 MB/2 MB/4 MB/	
	8 MB/16 MB block-level MD5 hash codes	

Figure 2: CDF (cumulative distribution function) of 1) original file size and 2) compressed file size, corresponding to our collected trace. For original files, the maximum size is 2.0 GB, the average size is 962 KB, and the median size is 7.5 KB. For compressed files, the maximum size is 1.97 GB, the average size is 732 KB, and the median size is 3.2 KB. Clearly, the tracked files can be effectively compressed on the whole, and the majority of them are small in size.

3. METHODOLOGY

This section describes our methodology for studying the *TUE* of cloud storage services. First, we introduce a real-world cloud storage trace collected to characterize the key impact factors. Next, we design a variety of benchmark experiments to uncover the key design choices of data sync mechanisms. Last but not the least, we provide an overview of the research results.

3.1 Real-world Cloud Storage Trace

Our measurement study takes advantage of a real-world user trace of cloud storage services. It is collected in several universities and companies in the US and China from Jul. 2013 to Mar. 2014, including 153 long-term users with 222,632 files inside their sync folders. Refer to Table 2 for the per service statistics. This trace is used to characterize the key impact factors with regard to the six widely used services. It is also used to guide the design of benchmark experiments and enable further macro-level analysis (in particular, the *TUE* related optimization opportunities of cloud storage services).

This cloud storage trace records detailed information of every tracked file in multiple aspects. Table 3 lists the concrete file attributes recorded. Figure 2 depicts the distributions of original file size and compressed file size corresponding to the trace. We have made this trace publicly available to benefit other researchers. It can be downloaded via the following link:

http://www.greenorbs.org/people/lzh/public/traces.zip

3.2 Benchmark Experiments

To obtain an in-depth understanding of *TUE* and the key design choices of data sync mechanisms, we design a variety of benchmarks for performing comprehensive controlled experiments. The benchmarks span multiple commercial cloud storage services, involving diverse client machines locating at distinct locations and network environments.

Cloud storage services. Among today's dozens of commercial cloud storage services, our research focuses on the following six mainstream services: Google Drive, OneDrive, Dropbox, Box, Ubuntu One, and SugarSync, as they are either the most popular (in terms of user base) or the most representative (in terms of data sync mechanism). Other cloud storage services are also briefly discussed when necessary.

Client locations. Since the above cloud storage services are mainly deployed in the US, we select two distinct locations to perform each experiment: MN (*i.e.,* Minnesota, US) and BJ (*i.e.,* Beijing, China). In a *coarse-grained* manner, MN represents a location close to the cloud: the bandwidth is nearly 20 Mbps and the latency $\in (42, 77)$ msec, while BJ represents a location remote from the cloud: the bandwidth is nearly 1.6 Mbps the latency $\in (200, 480)$ msec.

Controlled bandwidth and latency. To tune the network environment in a *fine-grained* manner, we interpose a pair of packet filters in the communication channel between the client and the cloud in MN. These filters enable fine-grained adjustment of the bandwidth (the maximum possible speed is 20 Mbps) and latency in either direction. Specifically, the packet filters are interposed by using an intermediate proxy that runs the Linux Netfilter/Iptables tool, thus behaving like a common software firewall.

Controlled file operations. We synthetically generate almost all kinds of file operations appearing in the literature. Moreover, these operations are applied upon both compressed and compressible files. These controlled file operations will be elaborated in § 4, § 5, and § 6.

Client machine hardware. A total of eight client machines are employed in the experiments: four in MN (*i.e.,* M1, M2, M3, and M4) and four in BJ (*i.e.,* B1, B2, B3, and B4). Their detailed hardware information is listed in Table 4. M1/B1 represents a typical client machine at the moment, M2/B2 an outdated machine, M3/B3 an advanced machine with SSD storage, and M4/B4 an Android smartphone. M1–M3 and B1–B3 are installed with Windows 7-SP1 and the Chrome-30.0 web browser.

Benchmark software and access methods. For each cloud storage service, all the experiments regarding M1–M3 and B1–B3 are performed with the latest version (as of Jan. 2014) of the client software on Windows 7. For the M4 and B4 smartphones, we experiment with the latest-version Android apps (as of Jan. 2014). The corresponding sync traffic (*i.e.,* incoming/outgoing packets) are recorded using Wireshark. For the Android smartphones, we route the network traffic through a PC that promiscuously monitors the packets using Wireshark.

Table 4: Hardware information of the experimental client machines.

Machine	CPU	Memory	Disk Storage
M1 @ MN	Quad-core Intel i5 @ 1.70 GHz	4 GB	7200 RPM, 500 GB
M2 @ MN	Intel Atom @ 1.00 GHz	1 GB	5400 RPM, 320 GB
M3 @ MN	Quad-core Intel i7 @ 1.90 GHz	4 GB	SSD, 250 GB
M4 @ MN	Dual-core ARM @ 1.50 GHz	1 GB	MicroSD, 16 GB
B1 @ BJ	Quad-core Intel i5 @ 1.70 GHz	4 GB	7200 RPM, 500 GB
B2 @ BJ	Intel Atom @ 1.00 GHz	1 GB	5400 RPM, 250 GB
B3 @ BJ	Quad-core Intel i7 @ 1.90 GHz	4 GB	SSD, 250 GB
B4 @ BJ	Dual-core ARM @ 1.53 GHz	1 GB	MicroSD, 16 GB

Table 5: Our major findings, their implications, and locations of relevant sections.

Simple File Operations	Implications
Section 4.1 (File creation): The majority (77%) of files in our trace are small in size (<100 KB), which may result in poor *TUE*.	For providers, nearly two thirds (66%) of small files can be logically combined into larger files for *batched data sync* (BDS). However, only Dropbox and Ubuntu One have partially implemented BDS so far.
Section 4.2 (File deletion): Deletion of a file usually incurs negligible sync traffic.	For users, no need to worry about the traffic for file deletion.
Section 4.3 (File modification): The majority (84%) of files are modified by users at least once. Most cloud storage services employ *full-file sync*, while Dropbox and SugarSync utilize *incremental data sync* (IDS) to save traffic for PC clients (but not for mobile or web-based access methods).	Most of today's cloud storage services are built on top of RESTful infrastructure (*e.g.,* Amazon S3, Microsoft Azure, and OpenStack Swift) that only support data access operations at the *full-file* level. *TUE* can be significantly improved by implementing IDS with an extra mid-layer that transforms MODIFY into GET + PUT + DELETE file operations.

Compression and Deduplication	Implications
Section 5.1 (Data compression): 52% of files can be effectively compressed. However, Google Drive, OneDrive, Box, and SugarSync never compress data, while Dropbox is the only one that compresses data for every access method.	For providers, data compression is able to reduce 24% of the total sync traffic. For users, PC clients are more likely to support compression versus mobile or web-based access methods.
Section 5.2 (Data deduplication): Although we observe that 18% of files can be deduplicated, most cloud storage services do not support data deduplication, especially for the web-based access method.	For providers, implementing compression and block-level deduplication together is technically challenging. Based on the trace analysis, we suggest providers implement compression and full-file deduplication since the two techniques work together seamlessly.

Frequent File Modifications	Implications
Section 6.1 (Sync deferment): Frequent modifications to a file often lead to large *TUE*. Some services deal with this issue by batching file updates using a fixed sync deferment. However, we find that fixed sync deferments are inefficient in some scenarios.	For providers, we demonstrate that an *adaptive sync defer* (ASD) mechanism that dynamically adjusts the sync deferment is superior to fixed sync deferment.
Section 6.2 (Network and hardware): Suprisingly, we observe that users with relatively low bandwidth, high latency, or slow hardware save on sync traffic, because their file updates are naturally batched together.	For users, in the presence of frequent file modifications, today's cloud storage services actually bring good news (in terms of *TUE*) to those users with relatively "poor" hardware or Internet access.

3.3 Overview of Our Major Findings

Based on the above methodology, we are able to thoroughly unravel the *TUE* relevant characteristics, design tradeoffs, and optimization opportunities of the six mainstream cloud storage services. The detailed research results (from simple to complex) will be presented in the following three sections: simple file operations (§ 4), compression and deduplication (§ 5), and frequent file modifications (§ 6). As an overview and a roadmap of our research results, Table 5 summarizes the major findings and their implications.

4. SIMPLE FILE OPERATIONS

This section presents our major measurement results, findings, and implications on the *TUE* of simple file operations. For each measurement, we first introduce the experiment process and results, and then unravel several interesting findings and implications. In this section, we do not mention the client locations, network environments, and hardware configurations, because the *TUE* of simple file operations is independent to these impact factors.

4.1 File Creation

[Experiment 1] : We first study the simple case of creating a highly compressed file of Z bytes inside the sync folder (we will further study the data compression in detail in § 5.1). Thereby, calculating the *TUE* of file creation becomes straightforward (*i.e., TUE* $= \frac{\text{Total sync traffic}}{Z \text{ bytes}}$). According to Figure 2, most compressed files are small in size (several KBs), and the maximum compressed file size is below 2.0 GB. Therefore, we experiment with

$Z \in \{1, 1\ K, 10\ K, 100\ K, 1\ M, 10\ M, 100\ M, 1\ G\}$.

Table 6: Sync traffic of a (compressed) file creation.

Service	PC client sync traffic (Bytes)				Web-based sync traffic (Bytes)				Mobile app sync traffic (Bytes)			
	1	1 K	1 M	10 M	1	1 K	1 M	10 M	1	1 K	1 M	10 M
Google Drive	9 K	10 K	1.13 M	11.2 M	6 K	7 K	1.06 M	10.6 M	32 K	71 K	1.27 M	11.0 M
OneDrive	19 K	20 K	1.14 M	11.4 M	28 K	31 K	1.11 M	11.7 M	29 K	44 K	1.23 M	10.7 M
Dropbox	38 K	40 K	1.28 M	12.5 M	31 K	37 K	1.09 M	10.6 M	18 K	32 K	1.08 M	10.9 M
Box	55 K	47 K	1.10 M	10.6 M	55 K	58 K	1.10 M	10.5 M	16 K	34 K	1.29 M	10.8 M
Ubuntu One	2 K	3 K	1.11 M	11.2 M	37 K	39 K	1.20 M	11.3 M	20 K	24 K	1.08 M	10.9 M
SugarSync	9 K	19 K	1.17 M	11.4 M	31 K	32 K	1.10 M	10.7 M	31 K	47 K	1.22 M	10.9 M

Table 7: Total traffic for synchronizing 100 compressed file creations. Each file is 1 KB in size.

Service	PC client		Web-based		Mobile app	
	Sync traffic	(*TUE*)	Sync traffic	(*TUE*)	Sync traffic	(*TUE*)
Google Drive	1.1 MB	(11)	1.2 MB	(12)	5.6 MB	(56)
OneDrive	1.3 MB	(13)	2.2 MB	(22)	1.9 MB	(19)
Dropbox	**120 KB**	**(1.2)**	*600 KB*	*6.0)*	*360 KB*	*(3.6)*
Box	1.2 MB	(12)	3.2 MB	(32)	3.2 MB	(32)
Ubuntu One	**140 KB**	**(1.4)**	*500 KB*	*(5.0)*	2.5 MB	(25)
SugarSync	0.9 MB	(9)	4.0 MB	(40)	1.5 MB	(15)

The second goal of **Experiment 1** is to get a quantitative understanding of the *overhead traffic*, as *TUE* heavily depends on the ratio of the overhead traffic over the total sync traffic. Synchronizing a file to the cloud always involves a certain amount of overhead traffic, which arises from TCP/HTTP(S) connection setup and maintenance, metadata delivery, etc. Specifically, the overhead traffic is equal to the total sync traffic excluding the payload traffic for delivering the file content, so in **Experiment 1**,

Overhead traffic ≈ Total sync traffic - Z bytes.

Table 6 lists the results of **Experiment 1** regarding the six concerned cloud storage services. We vary the file size from 1 B to 1 GB, but for brevity only list four typical sizes: 1 B, 1 KB, 1 MB, and 10 MB. The table records the sync traffic generated by the three typical service access methods: PC client, web (browser) based, and mobile app. In general, from Table 6 we have the following finding and implication:

- *TUE* for synchronizing a (compressed) file creation mainly depends on the file size. A *small* file results in big *TUE* up to 40000, while a *big* file incurs small *TUE* approaching 1.0. Therefore, for providers, a number of small files can be logically combined into a *moderate-size* file for *batched data sync* (BDS) to save traffic, in particular the overhead traffic.

This finding poses a key question: *What is a small size and what is a moderate size?* By plotting the *TUE vs.* File Size relationship (for PC clients) in Figure 3, we get an intuitive conclusion that a moderate size should be at least 100 KB and had better exceed 1 MB, in order to achieve small *TUE* – at most 1.5 and had better stay below 1.2. Here we only draw the curve for PC clients since the corresponding curves for web-based and mobile apps are similar.

As a consequence, small size is regarded as less than 100 KB, which together with Figure 2 reveals that the majority (77%) of tracked files are small in size (meanwhile, 81% in terms of compressed size). More importantly, by analyzing our collected trace, we find that nearly two-thirds (66%) of these small files can be created in batches and thus can effectively benefit from BDS.

[**Experiment 1'**] : Given that the BDS mechanism can effectively optimize *TUE*, a new question comes out: *Is BDS adopted by the*


**Figure 3: *TUE vs.* Size of the created file.**

six mainstream cloud storage services? To get the answer, we first generate 100 (distinct) highly compressed files, and then move all of them into the sync folder in a batch. Each file is 1 KB in size, so $TUE = \frac{\text{Total sync traffic}}{100 \text{ KB}}$. If BDS is adopted, the total sync traffic should be around 100 KB and *TUE* should be close to 1.0.

The results of **Experiment 1'** listed in Table 7 reveal that Dropbox and Ubuntu One have adopted BDS for PC clients. Further, it is possible that Dropbox has adopted BDS for web-based and mobile access methods, because the corresponding sync traffic (600 KB and 360 KB) is within an order of magnitude of the data update size (100 KB). Also, Ubuntu One may have used BDS in its web-based data synchronization, since the sync traffic (500 KB) lies between 600 KB and 360 KB. On the contrary, Google Drive, OneDrive, Box, and SugarSync have not adopted BDS yet.

4.2 File Deletion

[**Experiment 2**] : Each file created in **Experiment 1** is deleted after it is completely synchronized to the cloud, so as to acquire the sync traffic information of a file deletion.

The **Experiment 2** results indicate that deletion of a file usually generates negligible (< 100 KB) sync traffic, regardless of the cloud storage service, file size, or access method. The reason is straightforward: when a file f is deleted in the user's local sync fold-

Figure 4: Sync traffic of a random byte modification, corresponding to the three typical service access methods: (a) PC client, (b) Web-based, and (c) Mobile app. By comparing the three subfigures, we discover that only the PC clients of Dropbox and SugarSync utilize the incremental data sync (IDS) mechanism for improved network-level efficiency.

er, the user client just notifies the cloud to change some attributes of f rather than remove the content of f. In fact, such "fake deletion" also facilitates users' data recovery, such as the version rollback of a file. Naturally, we get the following implication:

- Cloud storage service users do not need to worry about the sync traffic when deleting a file.

4.3 File Modification and Sync Granularity

[Experiment 3] : The analysis of our collected cloud storage trace reveals that the majority (84%) of files are modified by users at least once. That is to say, file modifications are frequently made by cloud storage users. This subsection studies a simple case of file modification, *i.e.*, modifying a random byte in a compressed file of Z bytes inside the sync folder. In this case, $TUE = \frac{\text{Total sync traffic}}{1 \text{ Byte}}$. Similar as § 4.1, we experiment with

$Z \in \{1, 1\,K, 10\,K, 100\,K, 1\,M, 10\,M, 100\,M, 1\,G\}$

and plot the sync traffic of four typical sizes: $Z = 1$ K, 10 K, 100 K, and 1 M in Figure 4.

Figure 4 shows that today's cloud storage services generally utilize two kinds of data sync granularity: 1) full-file and 2) chunk-level. Accordingly, their data sync mechanisms are classified into *full-file sync* and *incremental sync* as follows:

- **Full-file sync.** Google Drive is an example to use the full-file sync mechanism. When a random byte is modified in a Z-byte compressed file, the resulting sync traffic is almost the same as that of creating a new Z-byte compressed file. In other words, Google Drive deals with each file modification by simply uploading the full content of the modified file to the cloud and then deleting the old file. Consequently, Google Drive is more suitable for hosting media files (like photos, music, and videos) which are rarely modified by users. The full-file sync mechanism is also employed by OneDrive, Box, Ubuntu One, Amazon Cloud Drive, and some popular cloud storage services in China like Kuaipan, Kanbox, Baidu CloudDisk, and 115 CloudDisk.

- **Incremental sync** (or **delta sync**). Dropbox (PC client) is an example to use the incremental data sync (IDS) mechanism. When a random byte modification happens, the resulting sync traffic stays around 50 KB, regardless of the size of the modified file. According to the working principle of the incremental sync algorithm: rsync [16], once a random byte is changed in a file f, in most cases the whole data chunk that contains this byte must be delivered for synchronizing f. Therefore, the sync granularity (*i.e.,* the chunk size C) can be approximately estimated as

$$C \approx \text{Total sync traffic} - \text{Overhead traffic}.$$

From the **Experiment 1** results, we understand that the overhead traffic of synchronizing a one-byte file with the Dropbox PC client is nearly 40 KB. Therefore, the data sync granularity of Dropbox PC client is estimated as: $C \approx 50$ KB $- 40$ KB $= 10$ KB. This is further validated by the recommended default chunk size (*i.e.,* from 700 B to 16 KB) in the original rsync implementation [15]. Moreover, we find that SugarSync, IDriveSync, and 360 CloudDisk also utilize the IDS mechanism for their PC clients.

On the contrary, as depicted in Figure 4 (b) and 4 (c), web-based apps and mobile apps for all the six services still use the full-file sync mechanism, probably because IDS is hard to implement in JavaScript (for web-based apps) or due to energy concerns (for mobile apps). Specifically, JavaScript is the most widely used script language for the development of web-based apps (including cloud storage apps). Nevertheless, for security concerns, JavaScript is unable to directly invoke file-level system calls/APIs like open, close, read, write, stat, rsync, and gzip [11]. Instead, JavaScript can only access users' local files in an indirect and constrained manner, which is of less efficiency in terms of implementing IDS.

In summary, we have the following finding about simple file modification and the data sync granularity:

- When a file modification is synchronized to the cloud, *TUE* is mostly affected by the data sync granularity which varies significantly among different cloud storage services. Most services simply use full-file sync, but some services (like Dropbox and SugarSync) utilize IDS to achieve improved network-level efficiency for PC clients.

Conflicts between IDS and RESTful infrastructure. Although enabling the IDS mechanism can help service providers reduce the data sync traffic, implementing IDS is not an easy job in practice. Most of today's cloud storage services (*e.g.,* OneDrive, Dropbox, and Ubuntu One) are built on top of RESTful infrastructure (*e.g.,* Amazon S3, Microsoft Azure, and OpenStack Swift). For simplifying both the providers' implementation complexities and the developers' programming complexities, RESTful infrastructure typically only supports data access operations at the *full-file* level, like PUT (upload a new file), GET (download a whole file), DELETE (delete a whole file), and so forth. Note that the PUT, GET, and DELETE operations may have aliases in other RESTful infrastructure.

Thus, enabling IDS usually requires an extra mid-layer for transforming MODIFY into GET + PUT + DELETE file operations in an *efficient* manner (like what Dropbox has done [25, 36]) [4]. Since file modifications frequently happen, implementing such a mid-layer is worthwhile for improved network-level efficiency.

---

[4]An alternative to enable IDS is to store every chunk of a file as a separate data object. When a file is modified, the modified chunks are deleted with the new chunks being stored as new objects; also, file metadata has to be updated (as what is used in Cumulus [43]).

Table 8: Sync traffic of a 10-MB text file creation. UP: The user uploads the file to the cloud. DN: The user downloads the file from the cloud.

Service	Sync traffic (MB)					
	PC client		Web-based		Mobile app	
	UP	DN	UP	DN	UP	DN
Google Drive	11.3	11.0	10.6	11.7	11.8	10.8
OneDrive	11.4	11.2	11.0	11.0	12.2	10.7
Dropbox	**6.1**	**5.5**	10.6	**5.5**	**8.1**	**5.5**
Box	10.6	11.2	10.5	11.3	10.4	11.1
Ubuntu One	**5.6**	**5.3**	10.9	**5.3**	**8.6**	10.6
SugarSync	11.3	11.5	10.4	10.7	11.6	11.8

5. COMPRESSION AND DEDUPLICATION

In a real-world storage system, compression and deduplication are the two most commonly used techniques for saving space and traffic. This section makes a detailed study of cloud storage compression and deduplication, from both the system designers' and users' perspectives.

5.1 Data Compression Level

[Experiment 4] : To study whether data updates are compressed before they are synchronized to the cloud, we create an X-byte text file inside the sync folder. As a small file is hard to compress, we experiment with $X = 1$ M, 10 M, 100 M, and 1 G. Each text file is filled with random English words. If data compression is actually used, the resulting sync traffic should be much less than the original file size. Furthermore, after each text file is completely synchronized, we download it from the cloud with a PC client, a web browser, and a mobile app, respectively, so as to examine whether the cloud delivers data updates in a compressed form.

As a typical case, the **Experiment 4** results corresponding to a 10-MB text file are listed in Table 8. First, in the file upload (UP) phase, among the six mainstream cloud storage services, only Dropbox and Ubuntu One compress data with PC clients and mobile apps. No service ever compresses data with web browsers. Further, we observe that the 10-MB text file can be compressed to nearly 4.5 MB using the highest-level WinZip compression on our desktop. Thus, for Dropbox and Ubuntu One, the compression level with PC clients seems moderate, while the compression level with mobile apps is quite low. The motivation of such a difference is intuitive: to reduce the battery consumption of mobile devices caused by the computation-intensive data compressions.

Next, in the file download (DN) phase, only Dropbox and Ubuntu One compress data with PC clients and web browsers, and the compression level is higher than that in the file upload phase. For mobile apps, only Dropbox compresses data.

By analyzing our collected cloud storage trace, we find that 52% of files can be *effectively compressed*. Here "effectively compressed" implies that $\frac{\text{Compressed file size}}{\text{Original file size}} < 90\%$ when the highest-level WinZip compression is applied. As a result, the total *compression ratio* ($= \frac{\text{Size of data before compression}}{\text{Size of data after compression}}$) regarding all the files recorded in the trace reaches 1.31. In other words, data compression is able to reduce 24% of the data sync traffic (compared with no compression). These observations lead to the following findings and implications:

- Data compression provides obvious benefits by reducing the sync traffic, but is not supported by every cloud storage service. Google Drive, OneDrive, Box, and SugarSync never

---

Algorithm 1 : Iterative Self Duplication Algorithm

1: Set the lower bound: $L = 0$ bytes, and the upper bound: $U = +\infty$ bytes. Guess a deduplication block size: B_1;

2: **Step 1**:
3:    Generate a new compressed file f_1 of B_1 bytes;
4:    Upload f_1 to the cloud. When f_1 is completely synchronized to the cloud, record the total sync traffic: Tr_1;

5: **Step 2**:
6:    Generate another file f_2 by appending f_1 to itself, that is $f_2 = f_1 + f_1$ (the so-called "self duplication");
7:    Upload f_2 to the cloud. When f_2 is completely synchronized to the cloud, record the total sync traffic: Tr_2;

8: **Step 3**:
9: **if** $Tr_2 \ll Tr_1$ **and** Tr_2 is small ($\approx$ tens of KBs) **then**
10:    B_1 is actually the deduplication block size (B);
11:    **exit**;
12: **else** there are two cases
13:    **case 1**: $Tr_2 < 2B_1$ and Tr_2 is not small (implying that $B_1 > B$) **then**
14:    Set B_1 as the upper bound: $U \leftarrow B_1$, and decrease the guessing value of B_1: $B_1 \leftarrow \frac{L+U}{2}$;
15:    **case 2**: $Tr_2 > 2B_1$ (implying that $B_1 < B$) **then**
16:    Set B_1 as the lower bound: $L \leftarrow B_1$, and increase the guessing value of B_1: $B_1 \leftarrow \frac{L+U}{2}$;
17: **goto Step 1**;

---

compress data, while Dropbox is the only service that compresses data for every access method.

- Web browser typically *does not* compress a file when uploading it to the cloud storage, probably also due to the limitations of JavaScript or other web-based script languages (refer to § 4.3). Besides, using a mobile app is usually not as efficient as using a PC client.

5.2 Data Deduplication Granularity

[Experiment 5] : Data deduplication is another potential method to reduce data sync traffic, with the intuition that users often upload duplicated files with similar content. Inferring the deduplication granularity of a cloud storage service requires some efforts, especially when the deduplication block size B (bytes) is not a power of two (*i.e.*, $B \neq 2^n$, where n is a positive integer) [5]. To measure the deduplication granularity, we design and implement Algorithm 1 (named the "Iterative Self Duplication Algorithm"). It infers the deduplication granularity by iteratively duplicating and uploading one or multiple synthetic file(s) and meanwhile analyzing the incurred data sync traffic. It is easy to prove that the iteration procedure can finish in $O(\log(B))$ rounds.

First, we study inter-file data deduplication with respect to an identical user account. By applying **Experiment 5** to the six mainstream cloud storage services, we figure out their data deduplication granularity in Table 9 (the 2nd column). In this table, "Full file" (only for Ubuntu One) means that data deduplication only happens at the full-file level, "4 MB" (only for Dropbox) indicates that the deduplication block size $B = 4$ MB, and "No" shows that there is no deduplication performed. Note that block-level deduplication naturally implies full-file deduplication, but *not* vice versa.

Second, we study cross-user data deduplication. For each cloud storage service, we first upload a file f to the cloud, and then

---

[5]The deduplication block size B (bytes) is traditionally a power of two [39], but we still have to thoughtfully consider the exception when $B \neq 2^n$.

Table 9: Data deduplication granularity. We do not list the web-based case because the web-based file synchronization typically *does not* apply data deduplication.

Service	Same user	Cross users
	PC client & Mobile app	PC client & Mobile app
Google Drive	No	No
OneDrive	No	No
Dropbox	4 MB	No
Box	No	No
Ubuntu One	Full file	Full file
SugarSync	No	No

Figure 5: *Deduplication ratio* (cross-user) *vs*. Block size. Here the deduplication ratio $= \frac{\text{Size of data before deduplication}}{\text{Size of data after deduplication}}$.

use another user account to upload f to the cloud again. In this case, the sync traffic should be trivial if full-file deduplication is performed across users. If the cross-user full-file deduplication is confirmed, **Experiment 5** is run again to figure out the accurate cross-user deduplication granularity; otherwise, we can conclude that there is no cross-user data deduplication at all. The results are also shown in Table 9 (the 3rd column). Obviously, only Dropbox employs a different cross-user data deduplication granularity from the identical-user case.

From the above measurements, we get the following findings and implications:

- A cloud storage service usually adopts the same data deduplication granularity for PC clients and mobile apps, while the web-based data synchronization typically *does not* apply data deduplication.

- By analyzing our collected trace, we find that cross-user data (block) duplication pervasively exists: even the full-file level *duplication ratio* ($= \frac{\text{Size of duplicate files}}{\text{Size of all files}}$) reaches 18.8%. However, most cloud storage services do not support cross-user deduplication (perhaps for privacy and security concerns) or block-level deduplication at the moment, thus losing considerable opportunities for optimizing *TUE*.

Further, we compare the two types of deduplication granularity to answer a question: *Is the block-level deduplication much better (i.e., has a much larger deduplication ratio) than the full-file deduplication?* Since the computation complexity of block-level deduplication is much higher than that of full-file deduplication, the answer could help decide whether or not the block-level deduplication is worthwhile. Note that when referring to the "file blocks", we are dividing files to blocks *in a simple and natural way*, that is to say, by starting from the head of a file with a fixed block size. So clearly, we are not dividing files to blocks in the best possible manner [19, 39] which is much more complicated and computation intensive.

As our collected trace contains both the full-file hash codes and the block-level (128 KB – 16 MB blocks) hash codes of each tracked file (refer to § 3.1, Table 3), we perform the trace-driven simulation to figure out the (cross-user) deduplication ratio when each deduplication granularity is adopted. The simulation results demonstrate that the block-level deduplication usually exhibits *trivial superiority* to the full-file deduplication, as shown in Figure 5. Therefore, we have the following implication:

- For providers, in terms of deduplication granularity, supporting full-file deduplication is basically sufficient.

Conflicts between compression and block-level deduplication. Although data deduplication can reduce the sync traffic, we notice that it has a potential performance conflict with data compression. Implementing block-level deduplication and compression together is technically challenging.

For cloud storage service providers, though storing and delivering data in its compressed form can effectively save storage space and sync traffic, it may significantly increase the (computation and I/O) complexity of block-level deduplication. Specifically, after a file (f) is delivered to the cloud storage in its compressed form (f'), f' must be first uncompressed to calculate each block's fingerprint, so as to enable block-level deduplication. Then, the uncompressed file must be deleted from disk. Furthermore, the above operations must be re-executed (in part) as long as one block of f is modified. It is basically unwise for a service provider to shift these operations to its user clients, unless the service provider does not care about user experience.

In this subsection we have known that block-level deduplication exhibits trivial superiority to full-file deduplication. Meanwhile, full-file deduplication is not challenged by data compression, because full-file deduplication can be directly performed on compressed files. Therefore, we suggest that providers implement full-file deduplication and compression since these two techniques work together seamlessly.
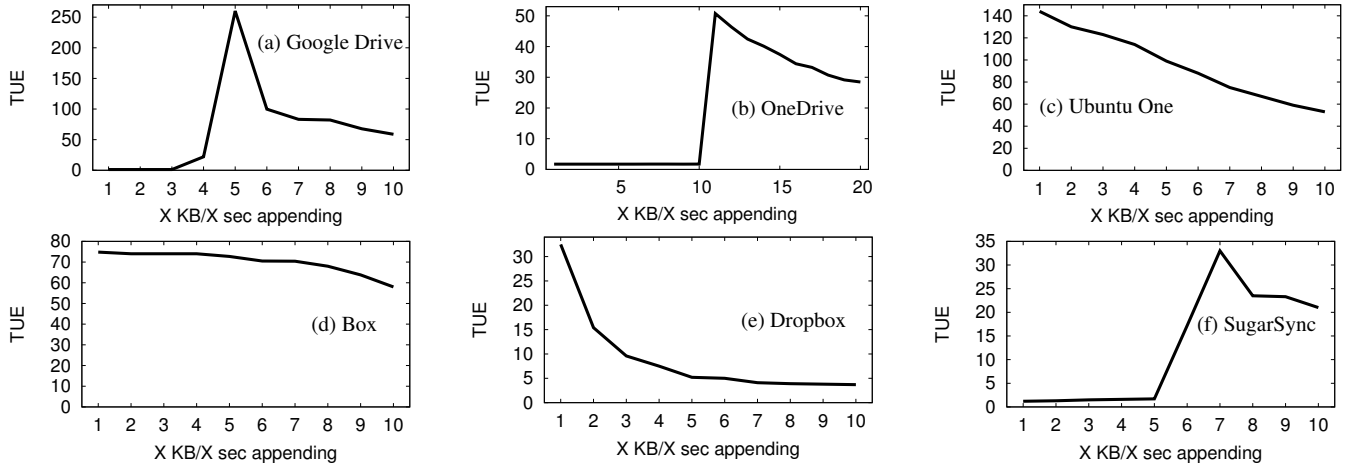
6. FREQUENT FILE MODIFICATIONS

In addition to backing up and retrieving files, cloud storage services are also widely used for collaboration, such as collaborative document editing, team project building, and database hosting. All the above mentioned advanced functions involve a special kind of file operations: *frequent modifications* to a file.

In § 4 and § 5, we have studied various simple file operations that are each performed at once. On the contrary, frequent modifications imply that a file is modified in a frequent and incremental manner. Thus, they exhibit diverse data update patterns in terms of data update *size* and *rate*. The large-scale trace collected by Drago *et al.* [12] reveals that for 8.5% of Dropbox users, more than 10% of their sync traffic is caused by frequent modifications [36]. Further, frequent modifications may well incur abundant overhead traffic that far exceeds the amount of useful data update traffic sent by the user client over time, which is referred to as the *traffic overuse problem* [36]. Besides, in this section we will elaborate on client locations, network environments, and hardware configurations, because the *TUE* of frequent file modifications is largely influenced by these factors.

6.1 Sync Deferment

[Experiment 6] : To experiment with frequent file modifications, we append X random kilobytes to an empty file inside the sync folder every X seconds, until the total appended bytes reach a certain size C (typically $C = 1$ MB). This is denoted as the "X KB/X

Figure 6: *TUE* **of the six cloud storage services in response to controlled frequent file modifications. Note that the subfigures have distinct Y-axes, and the (b) OneDrive subfigure has a different X-axis.**

sec" appending experiment. We use random bytes since they are difficult to compress, thus preventing file compression from influencing our measurements of *TUE*.

Our goal is three folds by doing this experiment: 1) We observe and understand the sync traffic and *TUE* in response to frequent modifications; 2) We aim to discover whether the cloud storage service has used the *sync deferment* in order to avoid or mitigate the traffic overuse problem; and 3) If the sync deferment is adopted, we want to measure how long the sync deferment is.

All the experiments in this section are performed using M1 @ MN (refer to Table 4) with 20 Mbps of bandwidth, and the latency (between M1 and each cloud) is between 42 msec and 77 msec. In terms of service access method, we only examine the PC client, because almost all the frequent modifications are generated from PC clients in practice [6]. Experiments with other benchmarks will be presented in § 6.2.

First, to investigate the impact of frequent file modifications on *TUE*, we examine the cases for $X \in \{1, 2, \cdots, 19, 20\}$. As depicted in Figure 6, the six mainstream cloud storage services exhibit diverse and interesting phenomena:

- Frequent modifications to a file often lead to large *TUE* (the aforementioned "traffic overuse problem"). As for the six mainstream services, the maximum *TUE* can reach 260, 51, 144, 75, 32, and 33, respectively.

- We observe (except in the sync defer cases: Figure 6 (a), 6 (b), and 6 (f)) that *TUE* generally decreases as the modification frequency $(= \frac{1024}{X})$ decreases. The reason is straightforward: though the total data update size is always $C = 1$ MB, a lower data update frequency implies fewer data sync events, and thus the overhead traffic is reduced.

A natural question is: *Why are the maximum* TUE *values of Google Drive (260), OneDrive (51), Ubuntu One (144), and Box (75) much larger than those of Dropbox (32) and SugarSync (33)?* The answer can be found from their data sync granularity (refer to § 4.3): Google Drive, OneDrive, Ubuntu One, and Box employ *full-file sync*, while Dropbox and SugarSync employ block-level

---

incremental sync which significantly improves the network-level traffic efficiency.

On the other hand, there do exist a few cases (in Figure 6 (a), 6 (b), and 6 (f)) where *TUE* is close to 1.0. According to our observations, (a) Google Drive, (b) OneDrive, and (f) SugarSync deal with the traffic overuse problem by batching file updates using a *fixed* sync deferment: T seconds (which cannot be re-configured by users). Figure 6 (a), 6 (b), and 6 (f) indicate that $T_{GoogleDrive} \in (3, 5)$ sec, $T_{SugarSync} \in (4, 6)$ sec, and $T_{OneDrive} \in (10, 11)$ sec. Moreover, to figure out a more accurate value of T, we further tune X from integers to floats. For example, we experiment with $X = 3.1, 3.2, \cdots, 4.9$ for $T_{GoogleDrive}$, and then find that $T_{GoogleDrive} \approx 4.2$ sec. Similarly, we find that $T_{SugarSync} \approx 6$ sec and $T_{OneDrive} \approx 10.5$ sec.

One may have the following question: *Is it possible that the deferred data synchronization of (a) Google Drive, (b) OneDrive, and (f) SugarSync is triggered by a* byte counter *or an* update counter *rather than the* time threshold (T)? In other words, the three concerned services may trigger the data synchronization once the number of uncommitted bytes or updates exceeds a certain value. This question can be addressed in two cases:

- Case 1: If the data synchronization is triggered by a *byte counter*, the resulting *TUE* would be close to 1.0 according to our previous study on the byte-counter based "efficient batched synchronization" (UDS) [36]. This is clearly not true as illustrated by Figure 6 (a), 6 (b), and 6 (f).

- Case 2: If the data synchronization is triggered by an *update counter*, the resulting *TUE* in Figure 6 (a), 6 (b), and 6 (f) would *linearly* decrease as the modification period (X sec) increases. Obviously, this is not true, either.

Therefore, we conclude that the deferred data synchronization is not triggered by a byte counter or an update counter.

Unfortunately, fixed sync deferments are limited in terms of usage scenarios. As shown in Figures 6 (a), 6 (b), and 6 (f), the traffic overuse problem still occurs when $X > T$.

To overcome the limitation of fixed sync deferments, we propose an *adaptive sync defer* (ASD) mechanism. ASD adaptively tunes its sync deferment (T_i) to follow the latest (say, the i-th) data update. In other words, when data updates happen more frequently, T_i gets shorter; when data updates happen less frequently, T_i gets

---

[6]The UIs (user interfaces) of web browser and mobile apps for cloud storage services are usually not fit for performing frequent modifications to a file.

Figure 7: *TUE* of (a) OneDrive, (b) Box, and (c) Dropbox on handling the "X KB/X sec" appending experiment, in Minnesota (MN) and Beijing (BJ), respectively.

Figure 8: *TUE* of Dropbox on handling the (a) "1 KB/sec" appending experiment with variable bandwidths, (b) "1 KB/sec" appending experiment with variable latencies, and (c) "X KB/X sec" appending experiment with distinct hardware configurations.

longer. In either case, T_i tends to be slightly longer than the latest inter-update time, so that frequent modifications can be properly batched for synchronization (without harming user experience). Specifically, T_i can be adapted in such an iterative manner:

$$T_i = \min\left(\frac{T_{i-1}}{2} + \frac{\Delta t_i}{2} + \epsilon,\ T_{max}\right) \quad (2)$$

where Δt_i is the inter-update time between the $(i\text{-}1)$-th and the i-th data updates, and $\epsilon \in (0, 1.0)$ is a small constant that guarantees T_i to be slightly longer than Δt_i in a small number of iteration rounds. T_{max} is also a constant representing the upper bound of T_i, as a too large T_i will harm user experience by bringing about intolerably long sync delay.

If Google Drive would utilize ASD on handling the "X KB/X sec" ($X > T_{GoogleDrive}$) appending experiments, the resulting *TUE* will be close to 1.0 rather than the original 260 ($X = 5$), 100 ($X = 6$), 83 ($X = 7$), and so forth. The situation is similar for OneDrive and SugarSync. More detailed performance evaluation of ASD can be found in our previous work [37].

6.2 Impact of Network and Hardware

In this subsection, we first study the impact of network and hardware on *TUE*, and then explore why they impact *TUE*.

[Experiment 7, Network environment] : To study the impact of network environment (including both bandwidth and latency) on *TUE*, we conduct the following two batches of experiments.

The first batch of experiments are performed on B1 @ BJ. It represents a relatively poor network environment: low bandwidth (nearly 1.6 Mbps) and long latency (between 200 msec and 480 msec) relative to the cloud, because the six mainstream cloud storage services are mainly deployed in US. After repeating **Experiments 1 – 6** in this network environment, we compare the results with the corresponding results by using M1 @ MN with abundant bandwidth (nearly 20 Mbps) and short latency (between 42 msec and 77 msec), which represents a good network environment.

The second batch of experiments are performed by using M1 @ MN with controlled bandwidth (between 1.6 Mbps and 20 Mbps) and latency (between 40 msec and 1000 msec), so that we are

able to get fine-grained results about how the network environment impacts *TUE*.

From the two batches of experiments, we mainly get the following findings and implications:

- *TUE* of a simple file operation is usually not affected by network environment.

- However, in the case of frequent file modifications, a user client with relatively low bandwidth or long latency can save more sync traffic.

Specifically, for the first batch of experiments, we plot the *TUE* of (a) OneDrive, (b) Box, and (c) Dropbox on handling the "X KB/X sec" appending experiment in Minnesota and Beijing in Figure 7 (a), 7 (b), and 7 (c), respectively. The situation of Google Drive and SugarSync is similar to Figure 7 (a), and the situation of Ubuntu One looks like Figure 7 (b). In each subfigure, the two curves ("@ MN" *vs.* "@ BJ") clearly illustrate that poor network environment leads to smaller *TUE*, especially when the modification period (X sec) is short (excluding the sync defer cases).

For the second batch of experiments, as a typical example, we plot the *TUE* of Dropbox on handling the "1 KB/sec" appending experiment with variable bandwidths and latencies in Figure 8 (a) and Figure 8 (b), respectively. In Figure 8 (a), the latency is fixed to around 50 msec and the bandwidth is tuned from 1.6 Mbps to 20 Mbps. In Figure 8 (b), the bandwidth is fixed to around 20 Mbps and the latency is tuned from 40 msec to 1000 msec. Obviously, higher bandwidth or shorter latency leads to larger TUE.

[Experiment 7', Hardware configuration] : Next, we examine the impact of hardware configuration on *TUE* by repeating **Experiments 1 – 6** with distinct client machines: M1 (a typical machine), M2 (an outdated machine), and M3 (an advanced machine). Their detailed hardware information is listed in Table 4. All the experiments are performed in Minnesota with abundant bandwidth (nearly 20 Mbps) and short latency (between 42 msec and 77 msec).

Through the **Experiment 7'** results, we observe that *TUE* of a simple file operation generally has no relation with hardware con-

figuration, but *TUE* of frequent file modifications is actually affected by hardware configuration. As a typical example, in Figure 8 (c) we plot the *TUE* of Dropbox on handling the "X KB/X sec" appending experiment with M1, M2, and M3. The three curves clearly demonstrate that slower hardware incurs less sync traffic.

Why do network environment and hardware configuration impact *TUE*? To explore the reason why network environment and hardware configuration impact *TUE*, we analyze the communication packets of data synchronization, in particular the TCP data flows. The analysis reveals that in the presence of frequent modifications to a file, the user client does *not* always synchronize every file modification to the cloud *separately*. Instead, the user client often batches multiple file modifications for data synchronization. Specifically, a new file modification (or a sequence of new file modifications) is synchronized to the cloud when at least the following two conditions are both satisfied:

- **Condition 1:** The previous file modification (or the previous batch of file modifications) has been completely synchronized to the cloud.

- **Condition 2:** The client machine has finished calculating the latest metatata of the modified file.

As to Condition 1, when the network environment is relatively poor, synchronizing the previous file modification (or the previous batch of file modifications) takes more time, so the client needs to wait for a longer period of time to synchronize the new file modification. As to Condition 2, when the client runs on top of slower hardware, calculating the latest metatata (which is computation-intensive) also requires a longer period of time. Because the failure of either condition will cause the new file modification (or the sequence of new file modifications) to be naturally batched, poor network environment or poor hardware increases the probability that a file modification gets batched, and thereby optimizes the *TUE*.

Finally, combining all the findings in this subsection, we get the following implication:

- In the case of frequent file modifications, today's cloud storage services actually bring good news (in terms of *TUE*) to those users with relatively poor hardware or Internet access.

7. DISCUSSION

While this paper mainly focuses on the *traffic costs* of cloud storage services, we keep in mind that the total costs of running a cloud storage service also involves the *computation costs*, *storage costs*, *operation costs*, and so forth. Therefore, we would like to further study and understand the traffic usage from an insider's point of view. In particular, we want to quantify the *tradeoff* between *TUE* and other system metrics. For example, regarding data sync granularity, incremental synchronization is a double-edge sword: It effectively saves traffic and storage compared with full-file synchronization, but it also puts more computational burden on both service providers and end users. Likewise, determining the best data compression level to achieve a good balance between traffic, storage, and computation deserves further research efforts.

Specifically, studying the aforementioned system tradeoffs would require at least the following three-fold information from cloud storage providers:

- *The user device composition* (*i.e.,* the percentages of PCs, tablets, and smartphones) is the most important information

required. For PCs, it is generally fine to sacrifice computation, storage, and/or network-level efficiency for better service quality (*e.g.,* faster synchronization of file operations). For example, PC clients usually maintain a local sync folder that stores a copy for almost every synchronized file at the cloud side. On the other hand, smartphones are sensitive to computation cost, storage space, and sometimes network overhead (in 2G/3G/4G modes). Accordingly, mobile apps usually maintain a small-size local folder that caches only a few most recently accessed files.

- *The logical interfaces of the storage infrastructure* decides the implementation difficulty and working efficiency of IDS (incremental data sync) for file modifications. The logical interfaces mainly include the RESTful (full-file level) interfaces, file-level interfaces, and block-level interfaces. For example, Microsoft Azure, Amazon S3, and OpenStack Swift provide RESTful interfaces, and thus implementing IDS on top of them is not an easy job. On the contrary, implementing IDS on top of a NFS-based infrastructure (with file-level interfaces) is quite straightforward. In addition, as GFS and HDFS provide seemingly file-level interfaces based on block-level infrastructure, the corresponding implementation difficulty and working efficiency of IDS lie between those with RESTful and file-level interfaces. Finally, the logical storage interfaces also impact the working efficiency of BDS (batched data sync) for small files.

- *The physical devices of the storage infrastructure* have non-negligible influence on the working efficiency and implementation (monetary) costs of a cloud storage service. Obviously, a single SSD is faster while much more expensive than a single HDD, but up to now it is still not clear which is the most cost-effective among an SSD cluster, an HDD cluster, a hybrid SSD+HDD cluster, or even a tape-based cluster [41]. Moreover, the performance of *the filesystem* can affect the working efficiency of a cloud storage service. For instance, OpenStack Swift works better with XFS than EXT3/EXT4, as pointed out by the official OpenStack development document [13]. In addition, even those seemingly independent infrastructure components may share deep, hidden dependencies that lead to unexpected *correlated failures*, thus undermining the redundancy efforts and working efficiencies of cloud storage services [46, 47].

We hope cloud storage service providers to release their proprietary traces to promote future research in this field.

8. RELATED WORK

As cloud storage services are becoming more pervasive and changing the way people store and share data, a number of research efforts have been made in academia, including the design and implementation of the service infrastructure [22, 43, 44, 31], integration services with various features and functionalities [29, 28, 23, 42, 36, 35, 34], performance measurement [33, 25, 24, 20, 45], as well as privacy and security issues [40, 21, 38, 27, 32]. While the previous work covers the data sync mechanism as one of the key operations and the resulting traffic usage, none of them tries to understand the *efficiency* of the traffic usage *comprehensively*. Due to the system complexity and implementation difference, one can hardly form a general and unified view of the traffic usage efficiency, not to mention the further improvement.

Our work is different from and complementary to previous studies by quantifying and optimizing traffic usage efficiency, the

pivotal network-level efficiency for cloud storage services. Based on the measurements and analysis of six state-of-the-art cloud storage services, we unravel the key impact factors and design choices that may significantly affect the traffic usage efficiency. Most importantly, we provide guidance and implications for both service providers and end users to economize their sync traffic usage.

Dropbox is one of the earliest and most popular cloud storage services, and its data sync mechanism has been studied in depth in [25, 36]. Through an ISP-level large-scale measurement, Drago *et al.* first uncover the performance bottlenecks of Dropbox due to both the system architecture and the data sync mechanism [25]. They suggest a bundling sync scheme with delayed sync ACK to improve the sync performance of Dropbox. In addition, Li *et al.* identify a pathological issue that may lead to the "traffic overuse problem" in Dropbox by uploading a large amount of unnecessary (overhead) traffic [36]. They propose an efficient batched sync algorithm (named UDS) to address this issue. Complementary to these studies, our results are not limited to Dropbox. Instead, we unravel the general factors that may significantly affect the data sync traffic. In consequence, our results are more general and applicable for designing network-level efficient cloud storage services, rather than improving one particular service.

Some measurement studies have partially covered the traffic usage of cloud storage services. Hu *et al.* examine "the good, the bad and the ugly" of four cloud storage services by comparing their traffic usage, delay time, and CPU usage of uploading new files [30]. They observe that the sync traffic usage varies substantially with factors such as file size, data compressibility, and data duplication levels. Drago *et al.* further compare the system capabilities of five cloud storage services, and find that each service has limitations with regard to data synchronization [24]. Both of these two studies confirm the importance of sync traffic usage, and the possibility of further optimizing the sync traffic usage.

In this paper, we zoom into the problem towards a comprehensive understanding of traffic usage efficiency. Different from the simplified benchmarks used in the above mentioned studies [36, 30, 24], we consider the diversity of access methods, client locations, hardware configurations, and network conditions to match the real-world usage. Indeed, we discover that these factors lead to different traffic usage patterns, some of which are even not expected. Last but not the least, different from previous studies that never consider mobile usage, one of our focus is the mobile usage of sync traffic – mobile users are those who mostly suffer from traffic overuse.

9. CONCLUSION AND FUTURE WORK

The tremendous increase in data sync traffic has brought growing pains to today's cloud storage services, in terms of both infrastructure support and monetary costs. Driven by this problem, this paper quantifies and analyzes the data sync traffic usage efficiency (*TUE*) of six widely used cloud storage services, using a real-world trace and comprehensive experiments. Our results and findings confirm that much of the data sync traffic is unnecessary and can be avoided or mitigated by careful design of data sync mechanisms. In other words, there is enormous space for optimizing the network-level efficiency of existing cloud storage services. We sincerely hope that our work can inspire the cloud storage designers to enhance their system and software, and meanwhile guide the users to pick appropriate services.

In Jun. 2014, Apple Inc. announced its cloud storage service, iCloud Drive, in its annual WWDC conference. iCloud Drive is planned as an important component of iCloud, the popular cloud service provided by Apple that has amassed over 300 million users. Therefore, in the near future, iCloud Drive may be ranked higher than several cloud storage services studied in this paper (*e.g.,* Dropbox and Box). To our knowledge, multiple groups of researchers (including us) have been anxious to investigate the network-level efficiency of iCloud Drive, in the hopes of harvesting novel and interesting discoveries. Unfortunately, up to now (Aug. 2014), the iCloud Drive service is still unavailable to common users [10]. In general, we believe that the network-level efficiency of iCloud Drive would be a promising research topic, as iCloud Drive lives in a unique and closed ecological system fully operated by Apple.

10. ACKNOWLEDGEMENTS

11. REFERENCES

[1] Amazon S3 pricing policy (Jan. 2014).
    http://aws.amazon.com/s3/#pricing.
[2] Bandwidth costs for cloud storage.
    http://blog.dshr.org/2012/11/bandwidth-costs-for-cloud-storage.html.
[3] Bandwidth limitations are a concern with cloud backup.
    http://searchdatabackup.techtarget.com/video/Bandwidth-limitations-are-a-concern-with-cloud-backup.
[4] Cisco Global Cloud Index: Forecast and Methodology, 2012-2017. Trend 3: Remote Data Services and Storage Access Services Growth.
    http://www.cisco.com/en/US/solutions/collateral/ns341/ns525/ns537/ns705/ns1175/Cloud_Index_White_Paper.html.
[5] Dirty Secrets: 5 Weaknesses of Cloud Storage Gateways.
    http://www.nasuni.com/blog/28-dirty_secrets_5_weaknesses_of_cloud_storage.
[6] Dropbox Is Now The Data Fabric Tying Together Devices For 100M Registered Users Who Save 1B Files A Day.
    http://techcrunch.com/2012/11/13/dropbox-100-million.
[7] Google Drive Now Has 10 Million Users: Available on iOS and Chrome OS. *http://techcrunch.com/2012/06/28/google-drive-now-has-10-million-users-available-on-ios-and-chrome-os-offline-editing-in-docs*.
[8] Hidden Costs of Cloud Storage.
    http://www.onlinefilestorage.com/hidden-costs-of-cloud-storage-1756.
[9] How fast is SkyDrive (OneDrive) growing?
    http://www.liveside.net/2012/10/27/how-fast-is-skydrive-growing.
[10] iCloud Drive features preview.
    http://www.apple.com/ios/ios8/icloud-drive.
[11] JavaScript Tutorials, Refernces, and Documentation.
    http://developer.mozilla.org/en-US/docs/Web/javascript.
[12] Large-scale Dropbox trace collected at the ISP level.
    http://traces.simpleweb.org/wiki/Dropbox_Traces.
[13] OpenStack Installation Guide for Ubuntu 12.04/14.04 (LTS).
    http://docs.openstack.org/icehouse/install-guide/install/apt/content.

[14] PUE (Power Usage Effectiveness). *http://en.wikipedia.org/wiki/Power_usage_effectiveness*.

[15] A question about the default chunk size of rsync. *http://lists.samba.org/archive/rsync/2001-November/000595.html*.

[16] rsync web site. *http://www.samba.org/rsync*.

[17] Why RESTful Design for Cloud is Best. *http://www.redhat.com/promo/summit/2010/presentations /cloud/fri/galder-945-why-RESTful/RestfulDesignJBWRH2010.pdf*.

[18] Wireshark network protocol analyzer. *http://www.wireshark.org*.

[19] B. Aggarwal, A. Akella, A. Anand, A. Balachandran, P. Chitnis, C. Muthukrishnan, R. Ramjee, and G. Varghese. EndRE: An End-system Redundancy Elimination Service for Enterprises. In *Proc. of NSDI*, pages 419–432. USENIX, 2010.

[20] A. Bergen, Y. Coady, and R. McGeer. Client Bandwidth: The Forgotten Metric of Online Storage Providers. In *Proc. of PacRim*, pages 543–548. IEEE, 2011.

[21] A. Bessani, M. Correia, B. Quaresma, F. André, and P. Sousa. DepSky: Dependable and Secure Storage in a Cloud-of-clouds. *ACM Transactions on Storage (TOS)*, 9(4):12, 2013.

[22] B. Calder, J. Wang, A. Ogus, N. Nilakantan, A. Skjolsvold, S. McKelvie, Y. Xu, et al. Windows Azure Storage: A Highly Available Cloud Storage Service with Strong Consistency. In *Proc. of SOSP*, pages 143–157. ACM, 2011.

[23] Y. Chen, K. Srinivasan, G. Goodson, and R. Katz. Design Implications for Enterprise Storage Systems via Multi-dimensional Trace Analysis. In *Proc. of SOSP*, pages 43–56. ACM, 2011.

[24] I. Drago, E. Bocchi, M. Mellia, H. Slatman, and A. Pras. Benchmarking Personal Cloud Storage. In *Proc. of IMC*, pages 205–212. ACM, 2013.

[25] I. Drago, M. Mellia, M.M Munafò, A. Sperotto, R. Sadre, and A. Pras. Inside Dropbox: Understanding Personal Cloud Storage Services. In *Proc. of IMC*, pages 481–494. ACM, 2012.

[26] R.T. Fielding. *Architectural Styles and the Design of Network-based Software Architectures*. PhD thesis, University of California, Irvine, 2000.

[27] S. Halevi, D. Harnik, B. Pinkas, and A. Shulman-Peleg. Proofs of Ownership in Remote Storage Systems. In *Proc. of CCS*, pages 491–500. ACM, 2011.

[28] D. Harnik, R. Kat, D. Sotnikov, A. Traeger, and O. Margalit. To Zip or Not to Zip: Effective Resource Usage for Real-Time Compression. In *Proc. of FAST*, pages 229–242. USENIX, 2013.

[29] D. Harnik, B. Pinkas, and A. Shulman-Peleg. Side Channels in Cloud Services: Deduplication in Cloud Storage. *IEEE Security & Privacy*, 8(6):40–47, 2010.

[30] W. Hu, T. Yang, and J.N. Matthews. The Good, the Bad and the Ugly of Consumer Cloud Storage. *ACM SIGOPS Operating Systems Review*, 44(3):110–115, 2010.

[31] Y. Huang, Z. Li, G. Liu, and Y. Dai. Cloud Download: Using Cloud Utilities to Achieve High-quality Content Distribution for Unpopular Videos. In *Proc. of ACM Multimedia*, pages 213–222. ACM, 2011.

[32] D. Kholia and P. Wegrzyn. Looking Inside the (Drop) box. In *Proc. of the 7th USENIX Workshop on Offensive Technologies (WOOT)*, 2013.

[33] A. Li, X. Yang, S. Kandula, and M. Zhang. CloudCmp: Comparing Public Cloud Providers. In *Proc. of IMC*, pages 1–14. ACM, 2010.

[34] Z. Li, Y. Huang, G. Liu, F. Wang, Z.-L. Zhang, and Y. Dai. Cloud Transcoder: Bridging the Format and Resolution Gap between Internet Videos and Mobile Devices. In *Proc. of NOSSDAV*, pages 33–38. ACM, 2012.

[35] Z. Li and J. Li. Deficiency of Scientific Research behind the Price War of Cloud Storage Services. *Communications of China Computer Federation (CCCF)*, 10(8):36–41, 2014.

[36] Z. Li, C. Wilson, Z. Jiang, Y. Liu, B.Y. Zhao, C. Jin, Z.-L. Zhang, and Y. Dai. Efficient Batched Synchronization in Dropbox-like Cloud Storage Services. In *Proc. of Middleware*, pages 307–327. Springer, 2013.

[37] Z. Li, Z.-L. Zhang, and Y. Dai. Coarse-grained Cloud Synchronization Mechanism Design May Lead to Severe Traffic Overuse. *Elsevier Journal of Tsinghua Science and Technology*, 18(3):286–297, 2013.

[38] P. Mahajan, S. Setty, S. Lee, A. Clement, L. Alvisi, M. Dahlin, and M. Walfish. Depot: Cloud Storage with Minimal Trust. *ACM Transactions on Computer Systems (TOCS)*, 29(4):12, 2011.

[39] D.T. Meyer and W.J. Bolosky. A Study of Practical Deduplication. *ACM Transactions on Storage (TOS)*, 7(4):14, 2012.

[40] M. Mulazzani, S. Schrittwieser, M. Leithner, M. Huber, and E. Weippl. Dark Clouds on the Horizon: Using Cloud Storage as Attack Vector and Online Slack Space. In *Proc. of USENIX Security*, 2011.

[41] V.S. Prakash, X. Zhao, Y. Wen, and W. Shi. Back to the Future: Using Magnetic Tapes in Cloud Based Storage Infrastructures. In *Proc. of Middleware*, pages 328–347. Springer, 2013.

[42] P. Shilane, M. Huang, G. Wallace, and W. Hsu. WAN-optimized Replication of Backup Datasets using Stream-informed Delta Compression. *ACM Transactions on Storage (TOS)*, 8(4):13, 2012.

[43] M. Vrable, S. Savage, and G.M. Voelker. Cumulus: Filesystem Backup to the Cloud. *ACM Transactions on Storage (TOS)*, 5(4):14, 2009.

[44] M. Vrable, S. Savage, and G.M. Voelker. Bluesky: A Cloud-backed File System for the Enterprise. In *Proc. of FAST*. USENIX, 2012.

[45] G. Wallace, F. Douglis, H. Qian, P. Shilane, S. Smaldone, M. Chamness, and W. Hsu. Characteristics of Backup Workloads in Production Systems. In *Proc. of FAST*. USENIX, 2012.

[46] E. Zhai, R. Chen, D.I. Wolinsky, and B. Ford. An Untold Story of Redundant Clouds: Making Your Service Deployment Truly Reliable. In *Proc. of HotDep*. ACM, 2013.

[47] E. Zhai, R. Chen, D.I. Wolinsky, and B. Ford. Heading Off Correlated Failures through Independence-as-a-Service. In *Proc. of OSDI*. USENIX, 2014.

[48] Y. Zhang, C. Dragga, A. Arpaci-Dusseau, and R. Arpaci-Dusseau. ViewBox: Integrating Local File Systems with Cloud Storage Services. In *Proc. of FAST*, pages 119–132. USENIX, 2014.